

Harsh Gupta

☎ +91 87006 10534 | ✉ harshgupta15march@gmail.com | [in linkedin.com/in/harshgupta](https://www.linkedin.com/in/harshgupta) | github.com/Harshgt57 | [gfg/harshgt57](https://gfg.harshgt57)

PROFESSIONAL SUMMARY

Data Engineer with 3+ years of experience designing, optimizing, and migrating enterprise ETL/ELT pipelines using **Azure Databricks, Apache Spark (PySpark/Spark SQL), Azure Synapse, and Delta Lake**. Proven track record of migrating legacy SAP HANA and on-premise workloads to the cloud, cutting data latency by up to 3 hours, and accelerating processing efficiency by 50–60%. Experienced across the full data lifecycle—from incremental data ingestion and Medallion architectures to high-performance BI reporting and automated pipeline monitoring. Certified in Microsoft Fabric, Azure, and Databricks.

SKILLS SUMMARY

- **Programming & Querying:** Python, PySpark, Spark SQL, SQL (Complex Joins, Window Functions, CTEs)
- **Big Data & Cloud Platforms:** Azure Databricks, Apache Spark, Delta Lake, Microsoft Fabric, Azure Synapse Analytics
- **Data Integration & Storage:** Azure Data Factory (ADF), Azure Data Lake Storage (ADLS Gen2), Databricks Workflows
- **Data Engineering Concepts:** Medallion Architecture (Bronze/Silver/Gold), ETL/ELT, SCD Type 1/2, Partitioning, AQE, Broadcast Joins, Data Validation & Reconciliation
- **Others:** Azure DevOps, Git, CI/CD Pipelines, Power BI, Tableau, Pipeline Logging & Performance Monitoring, SSMS

EXPERIENCE

- **Celebal Technologies** Noida, India
Senior Software Engineer (Promoted from Software Engineer) | (Databricks | Azure | Spark) May 2025 - Current
 - **Migrated SAP HANA calculation views and Hadoop-based workloads to Databricks**, engineered end-to-end scalable ETL/ELT data pipelines using **PySpark, Spark SQL, and Delta Lake**, leveraging **partitioning, broadcast joins, and Adaptive Query Execution (AQE)** across Sales, P&L, and Petchem business domains.
 - **Redesigned the Polyester reporting data model by consolidating 10+ tables and 15+ views supporting over 20 reports into a single precomputed table**, enabling a single refresh across reports, reducing maintenance and report-specific processing, and keeping report data synchronized.
 - **Improved Polyester, EBITDA, and Sales dashboard performance to 5–6 seconds** by shifting transformation and KPI computation from the Tableau/report layer to precomputed Databricks tables, reducing report-side processing and data-access latency.
 - **Migrated SAP BEx-based reports including MEG, FID, Chemical Status, DCA Invoice to Databricks**, implementing DBR-based ETL/transformation logic and performing HANA-to-Databricks reconciliation for data accuracy.
 - **Executed HANA-to-Databricks data reconciliation and troubleshooting**, resolving critical schema drift, join logic anomalies, and business-rule discrepancies across production datasets.
 - Architected **Medallion (Bronze/Silver/Gold) Lakehouse pipelines** on ADLS Gen2, integrating Delta MERGE (SCD Type 1/2) for automated historical backfills and incremental refreshes orchestrated via Databricks Workflows.
 - Mentored junior engineers and interns on **Databricks, PySpark, SQL, ETL, and debugging**, conducted knowledge-sharing sessions on **Spark/Databricks** and production data engineering practices.
- **MAQ Software** Noida, India
Software Engineer (Promoted from Associate Software Engineer) Jan 2023 - May 2025
 - Migrated data pipelines by eliminating an intermediate server, enabling real-time data availability and **reducing latency by 3 hours**.
 - Engineered and orchestrated cloud ETL pipelines leveraging **Azure Databricks, Azure Synapse and Apache Spark**, cutting data processing runtime by **50–60%**.
 - Developed and optimized **Power BI dashboards**, driving business decisions and contributing to **up to 65% cost reduction**.
 - Reduced report load time (PLT) by **70–80%**, significantly improving user experience and performance.
 - Designed and published **10+ reports**, reducing manual data analysis effort by **60%**.
 - Implemented logging and monitoring frameworks in Spark, enhancing **pipeline observability and troubleshooting**.
 - Collaborated on **QA, code reviews, and training**, improving team productivity and onboarding efficiency.

EDUCATION

- **G.L. Bajaj Institute Of Technology and Management** Greater Noida, India
B.Tech in Computer Science Engineering; CGPA: 8.47 Aug 2019 - Jun 2023

CERTIFICATION

Microsoft Certified: Fabric Analytics Engineer Associate , Fabric Data Engineer Associate , Azure Developer Associate 

Databricks Certified: Data Engineer Associate , Data Engineer Professional , Spark Developer Associate 